

# Simulation

## Lecture 8

# Quiz 2 Information

- 10 to 12 questions.
  - Lecture 5 to 8 but you may see some discrete distributions (in regards to simulation).
  - Multiple choice, multiple answer, calculation, and coding in R.
- We will release the practice quiz by Thursday.
- No multivariate calculus will be evaluated (i.e., partial derivatives and multiple integrals).

# A quick overview before our last lecture!

# How to deal with uncertainty?

- So far, we have seen **many quantities** that help us communicate an uncertain outcome:
  - PDF/PMF
  - Odds
  - Mode/Mean/Median
  - Entropy/Variance/standard deviation

## However...

- Sometimes, it is **difficult** to compute them.
- In these situations, we can use **simulation**.

# Outline

1. Review on Random Samples
2. Random seeds
3. Generating Random Samples
4. Running Simulations
5. Multi-Step Simulations

# 1. Review on Random Samples

- A **random sample** is a collection of random variables.
- A random sample of size  $n$ :  $X_1, \dots, X_n$ .
- Examples:
  - The outcomes of ten dice rolls ( $n = 10$ ).

# Assumption about Random Samples

- Unless we make additional sampling assumptions, a random sample is assumed to be independent and identically distributed (iid).



## 2. Seeds

- We use algorithms to generate pseudorandom numbers, which is not truly random.
- These numbers appear random, but are actually generated by a deterministic process.

# Stochastic vs Deterministic

- The word **stochastic** refers to having some uncertain outcome.
- The opposite of stochastic is **deterministic**: an outcome that will be known with 100% certainty (0 entropy).

# Example of a pseudorandom number generator

- Starting with a number  $x_0 \in (0, 1)$ , we generate the next number by:

$$x_{i+1} = 4x_i(1 - x_i).$$

- The generator produces pseudorandom numbers between 0 and 1 ( $X \sim \text{Uniform}(0, 1)$ ).


# The output

- Here is the resulting sequence when we start with  $x_0 = 0.3$  and iterate  $n = 1000$  times:

R Code [↺ Start Over](#) [▶ Run Code](#)

```
1 n <- 1000
2 x <- 0.3
3 for (i in 1:n) x[i + 1] <- 4 * x[i] * (1 - x[i])
4 head(x, 10)
```

# Plotting values from vector **x**





# All pseudorandom number generators have some pitfalls

- The generated sequence is **deterministic**, but it looks like a random sample.
- Moreover, in this example, neighbouring pairs **are not independent of each other**.
- Some sophisticated algorithms can produce outcomes that more closely resemble a random sample.



# The seed or random state

- The **seed** (or **random state**) in a pseudorandom number generator is some initial values that determines the generated sequence.
  - In this example,  $x_0$  is the initial value.
- In **R**, we can use `set.seed()`.
- In **Python**, we can use `numpy.random.seed()`.
- It is important to set seed to make sure our result is **reproducible**.



### 3. Generating Random Samples: Code

- We will look at some **R** and **Python** functions.
- Note we will focus on discrete distributions here.

## 3.1. Sampling from Finite Number of Outcomes

- **R**: `sample()`.
- **Python**: `numpy.random.choice()`.
- Both **R** and **Python** uses their own pseudorandom number generators.

# sample() in R

```
sample(x, size, replace = FALSE, prob = NULL)
```

- **x** is the vector of outcome.
- **size** is the desired sample size.
- **replace = TRUE** for sampling with replacement.
- **prob** is the vector of the probabilities of the outcomes respective to **x**.
- If **prob** does not add up to 1, **R** will automatically adjusts the probabilities so that they add up to 1.

## R Example

- Here is an example of generating  $n = 10$  items using the Mario Kart item distribution from [lecture1](#).

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 set.seed(1)
2 outcomes <- c("banana", "bob-omb", "coin", "horn", "shell")
3 probs <- c(0.12, 0.05, 0.75, 0.03, 0.05)
4 n <- 10
5 sample(outcomes, size = n, replace = TRUE, prob = probs)
```

# `numpy.random.choice()` in Python

```
random.choice(a, size=None, replace=True, p=None)
```

- `a` is the array of outcomes.
- `size` is the desired sample size.
- `p` is the array of the probabilities respective to `x`.
- `p` needs to add up to 1.

# Python Example

- Using the Mario Kart example again, we have the following:

Python Code ↺ Start Over ▶ Run Code

```
1 import numpy
2 numpy.random.seed(551)
3 outcomes = ["banana", "bob-omb", "coin", "horn", "shell"]
4 probs = [0.12, 0.05, 0.75, 0.03, 0.05]
5 n = 10
6 numpy.random.choice(a = outcomes, size = n, p = probs)
```

## 3.2. Sampling from a Distribution Family

- **R:** `r<dist>()` function, where `<dist>` is replaced with a short form of the distribution family's name.
- **Python:** `scipy.stats` from the `scipy` library.

The table below summarizes the functions for discrete distribution families

Family	R Function	Python Function
Binomial	<code>rbinom()</code>	<code>scipy.stats.binom.rvs()</code>
Geometric	<code>rgeom()</code>	<code>scipy.stats.geom.rvs()</code>
Negative Binomial	<code>rnbinom()</code>	<code>scipy.stats.nbinom.rvs()</code>
Poisson	<code>rpois()</code>	<code>scipy.stats.poisson.rvs()</code>



# How to use these functions?

```
rbinom(n, ...)
```

```
scipy.stats.binom.rvs(..., size=1)
```

- Sample size:
  - For **R**: the argument **n**, which comes first.
  - For **Python**: the argument **size**, which comes last.

# How to use these functions?

```
rbinom(n, size, prob)
```

```
scipy.stats.binom.rvs(n, p, size=1)
```

- In both languages, each parameter has its own argument.
- Sometimes, we have the option for different parameterizations.
- Be sure to specify the exact number of parameters required to identify the distribution!

# Generate $n = 10$ Binomial random numbers with $p = 0.6$ and 5 trial

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 set.seed(551)
2 rbinom_output <- rbinom(n = 10, size = 5, prob = 0.6)
3 rbinom_output
```

Python Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 import scipy
2 import numpy
3 numpy.random.seed(551) # scipy.stats uses numpy.random to generate its random numbers
4 scipy.stats.binom.rvs(n = 5, p = 0.6, size = 10)
```

# Negative Binomial Example

- Suppose the following:

$$X \sim \text{Negative Binomial}(k, p).$$

- $X$  refers to the number of failures in independent Bernoulli trials before experiencing  $k$  successes with probability  $p$ .

## R Function `rnbinom()`

- We can sample `n` Negative Binomial-distributed random numbers with `size` as  $k = 5$  and `prob` as  $p = 0.6$ :

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 set.seed(551)
2 rnbinom(n = 10, size = 5, prob = 0.6)
```

# Specifying too few parameters would result in an error

R Code ↺ Start Over ▶ Run Code

```
1 set.seed(551)
2 rbinom(n = 10, size = 5)
```

- R would indicate that `prob` is missing.

## Another way to use `rnbinom()` in R!

- Recall the expected value of a Negative Binomial random variable is

$$\mu = \mathbb{E}(X) = \frac{k(1 - p)}{p} = \frac{5(1 - 0.6)}{0.6} = 3.33.$$

- A Negative Binomial distribution can also be parameterized with  $k$  and its mean.

## Therefore...

- We can use the the argument `mu` in the random number generator `rnbinoom()`. Note `mu` refers to  $\mu = \mathbb{E}(X)$ :

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 set.seed(551)
2 rnbinoom(n = 10, size = 5, mu = 3.33)
```

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 set.seed(551)
2 rnbinoom(n = 10, size = 5, prob = 0.6)
```



# iClicker Question

Suppose you want to simulate hourly bank branch queues of customers. **Historically**, hourly queues show an average of 10 people.

What **R** random number generator will you use to simulate 20 random numbers?

- A. `rpois(n = 20, lambda = 1 / 10)`
- B. `rbinom(n = 20, size = 10, prob = 1 / 10)`
- C. `rpois(n = 20, lambda = 10)`
- D. `rgeom(n = 20, prob = 1 / 10)`

# Answer

- We need to simulate numbers of people in a time interval (i.e., hourly). These numbers are **integers** and non-negative.
- Therefore, we can go ahead with Poisson.
- Now, the parameter to use is the rate  $\lambda = 10$  people per hour.

# How do we read the function's output below?

- These 20 numbers indicate the simulated number of people per hour: 12 people, 10 people, 11 people, ...

R Code [↺ Start Over](#) [▶ Run Code](#)

```
1 set.seed(551) # Reproducibility
2 rpois(n = 20, lambda = 10)
```

# iClicker Question

Suppose you want to simulate the number of non-authentic bubble tea shops you will try before encountering your very first authentic one in Vancouver. Overall, **it is known that 70% of bubble tea shops in Vancouver are considered non-authentic.**

What **R** random number generator will you use to simulate 15 random numbers?

A. `rbinom(n = 15, size = 15, prob = 0.7)`

B. `rgeom(n = 15, prob = 0.3)`

C. `rbinom(n = 15, size = 15, prob = 0.3)`

B. `rgeom(n = 15, prob = 0.7)`

# Answer

- We need to simulate the number of failures (i.e., non-authentic) before encountering the first success (i.e., first authentic).
- Therefore, a Geometric distribution with  $p = 0.3$  is our way to go! Note that  $p$  is the probability of success (i.e., authentic).

## How do we read the function's output below?

- We have 15 random numbers indicating the simulated number of non-authentic places before encountering the first authentic one: 0 places, 0 places, 0 places, 5 places, ...

R Code [↺ Start Over](#) [▶ Run Code](#)

```
1 set.seed(551) # Reproducibility
2 rgeom(n = 15, prob = 0.3)
```

# About Student Experience of Instruction (SEI) Surveys

- They help us improve our teaching!
- UBC uses these in evaluating professors for hiring, and tenure promotion.

# What to comment on:

- Things that worked well for you!
- Things that could be improved!



# How to write your comments:

- Be constructive.
- Be respectful.
- Be concrete.

# SEI Surveys (10-15 mins break)



You voices can make a difference!

- Please go to this website: <https://seoi.ubc.ca/surveys> or scan the QR code, and fill out the survey for
  - Vincent Liu (**Instructor**)
  - Mahdi Asmae (**TA**)
  - Anne-Sophie Fratzscher (**TA**)
  - Cindy Zhang (**TA**)

## 4. Running Simulations

- There are two ways to calculate probabilistic quantities (e.g., means and probabilities):
  1. The **distribution-based approach** (using the distribution), resulting in **true values**.
  2. The **empirical approach** (using data), resulting in **approximate values** that improve as the sample size increases.

# Example: The Mean

- The mean of a discrete random variable  $X$  can be calculated as

$$\mathbb{E}(X) = \sum_x x \cdot P(X = x).$$

- Or can be approximated using the **empirical** approach from a random sample  $X_1, \dots, X_n$  by

$$\mathbb{E}(X) \approx \frac{1}{n} \sum_{i=1}^n X_i.$$

## 4.1. R functions for Calculating Empirical Quantities

- `mean()` calculates the sample average

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i.$$

- `var()` calculates the sample variance

$$S^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2.$$

## Furthermore...

- `sd()` calculates the sample standard deviation.
- `quantile()` calculates the **empirical  $p$ -quantile**: the  $np$ 'th largest (rounded up) observation in a sample of size  $n$ .
- To get the **entire PMF**, use the `table()` function, or more conveniently, the `janitor::tabyl()` function.
- To get the mode, either get it manually using the `table()` or `janitor::tabyl()` function, or you can use `DescTools::Mode()`.

## 4.2. Basic Simulation

- Consider a **random person** dating via a **dating app** with a probability of having a successful date of 0.7.



## Also...

- Suppose we want to evaluate the number of failed dates before they experience 5 successful dates.
- Define  $X$  to be the number of failed dates before experiencing 5 successful ones, then

$$X \sim \text{Negative Binomial}(k = 5, p = 0.7).$$



# Comparing distribution-based and empirical approach

- First, generate our random sample ( $n = 10000$  observations).

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 set.seed(551)
2 k <- 5
3 p <- 0.7
4 n <- 10000
5 random_sample <- rbinom(n, size = k, prob = p)
6 head(random_sample, 30) # Showing the first 30 random numbers in vector random_sample
```

## 4.2.1. Mean

- **Theoretically**, the mean of  $X$  is  $\mathbb{E}(X) = \frac{k(1-p)}{p}$ .

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 (k * (1 - p)) / p
```

- **Empirically**, we can approximate  $\mathbb{E}(X)$  with the sample average in `random_sample`:

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 mean(random_sample)
```

## 4.2.2. Variance

- **Theoretically**, the variance of  $X$  is  $\text{Var}(X) = \frac{k(1-p)}{p^2}$ .

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 (k * (1 - p)) / p^2
```

- **Empirically**, we can approximate  $\text{Var}(X)$  with the sample variance in `random_sample`:

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 var(random_sample)
```

### 4.2.3. Standard deviation

- **Theoretically**, the standard deviation of  $X$  is

$$\text{sd}(X) = \sqrt{\frac{k(1-p)}{p^2}}.$$

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 sqrt((1 - p) * k / p^2)
```

- **Empirically**, we can approximate  $\text{sd}(X)$  with the sample standard deviation in `random_sample`:

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 sd(random_sample)
```

## 4.2.4. Probability of Seeing 0 Failures (i.e., 0 failed dates!)

- **Theoretically**, this probability can be computed as

$$P(X = 0) = \binom{k-1}{0} p^k (1-p)^x.$$

- Using **R**, we can compute this probability directly:

R Code

↺ Start Over

▶ Run Code

```
1 dnbinom(x = 0, size = k, prob = p)
```

## Empirically...

- We can approximate  $P(X = 0)$  by counting the number of random numbers equal to 0 in `random_sample` and dividing this count by  $n = 10000$ .

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 head(random_sample) # First 6 random numbers out of 10000
2 head(random_sample == 0) # First 6 logical values out of 10000
3 mean(random_sample == 0) # Using function mean()
```

## 4.2.5. Probability Mass Function

- We can also do it for  $P(X = i)$  with  $i = 1, 2, \dots$  (i.e., **the whole PMF!**).

R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 library(tidyverse)
2 library(janitor)
3
4 PMF <- tabyl(random_sample) %>%
5   select(x = random_sample, Empirical = percent) %>%
6   mutate(Theoretical = dnbinom(x, size = k, prob = p))
7
8 PMF <- PMF %>%
9   mutate(
10     Theoretical = round(Theoretical, 4),
11     Empirical = round(Empirical, 4)
12   )
```

# Showing the output **PMF**


R Code ↺ Start Over ▶ Run Code

1 PMF

- Note that the probabilities are similar!



# Plotting both PMFs



## 4.2.6. Mode

- The mode is the outcome with the **largest probability**.
- From our previous plots, we can see that the mode is 1.

R Code ↺ Start Over ▶ Run Code

```
1 ## Theoretical
2 PMF %>% filter(Theoretical == max(Theoretical)) %>%
3   pull(x)
```

R Code ↺ Start Over ▶ Run Code

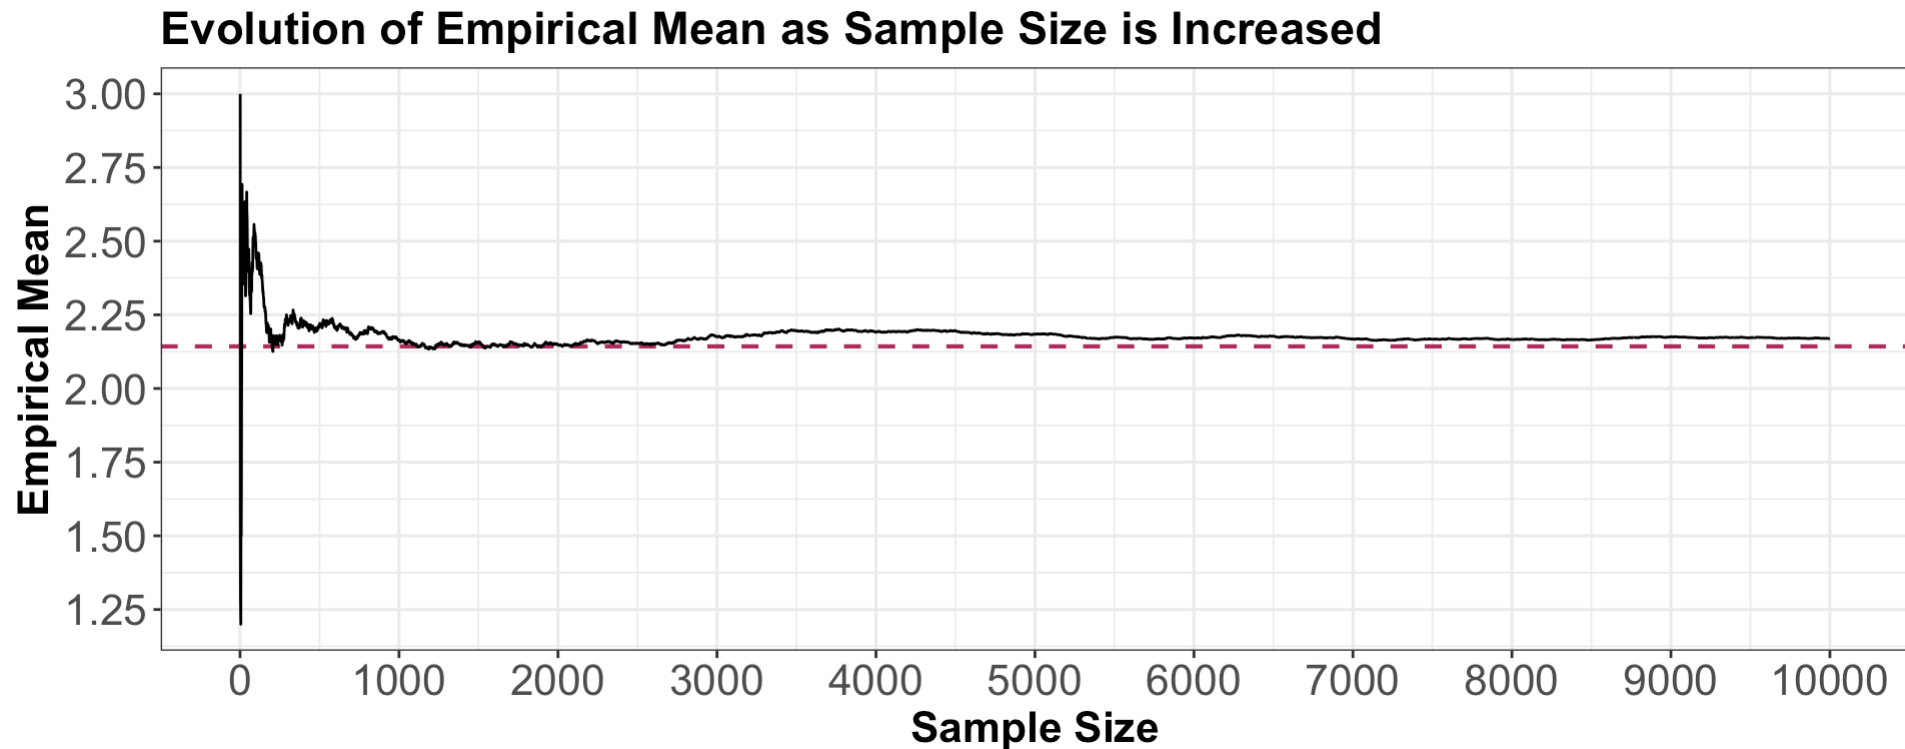
```
1 ## Empirical
2 PMF %>% filter(Empirical == max(Empirical)) %>%
3   pull(x)
```

# Why is our simulation so accurate?

- The Law of Large of Numbers states that, as we increase our sample size  $n$ , our empirical mean converges to the true mean.
- That is,  $\bar{X} \rightarrow \mu$  as  $n \rightarrow \infty$ .

# Increasing the sample size $n$ in our example

$X \sim \text{Negative Binomial}(k = 5, p = 0.7)$  with  $\mathbb{E}(X) = 2.14$ .



# 5. Multi-Step Simulations

- Simulation gets more interesting when we want to calculate things for a random variable that transforms and/or combines multiple random variables.

## For example...

- Consider a random variable  $T$  that we can obtain as follows:

$$X \sim \text{Poisson}(\lambda = 5)$$

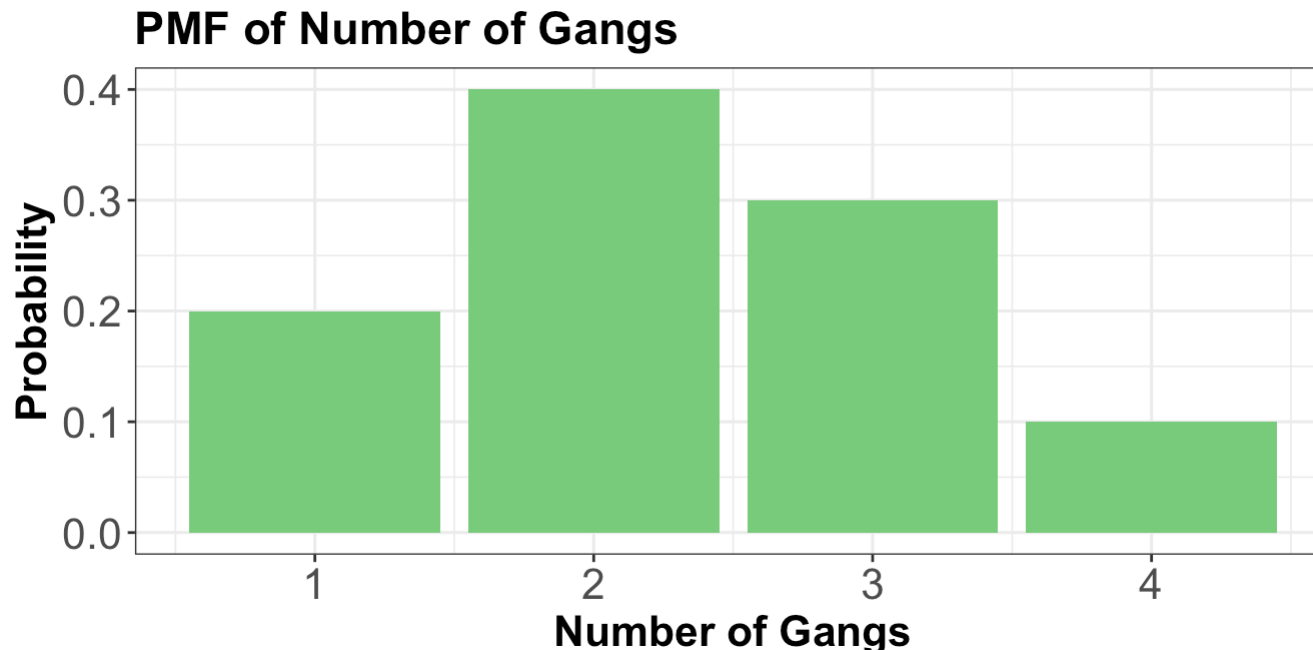
$$T = \sum_{i=1}^X D_i.$$

where each  $D_i$  are *iid* with some specified distribution.

- We would first generate  $X$ , then generate  $X$  values of  $D_i$ , then sum those up to get  $T$ .

# The Port of Vancouver Example

- Whenever a ship arrives, they request a certain number of **gangs** (groups of people) to help unload the ship.
- Let  $D_i$  denotes the random variable for the  $i$ -th ship with the following PMF:



# Simulation function

- Given a number of ships  $s$ , the function simulates the number of gangs requested, and sums up the gangs demand.
- That is, it simulates  $T = \sum_{i=1}^s D_i$ .

R Code

↺ Start Over

▶ Run Code

```

1 gang <- 1:4
2 p <- c(0.2, 0.4, 0.3, 0.1)
3
4 #' Generate gang demand
5 #'
6 #' Simulates the GRAND TOTAL number of gangs requested, if each ship
7 #' requests a random number of gangs.
8 #'
9 #' @param n_ships Number of ships that are making demands.
10 #' @param gangs Possible gang demands made by a ship.
11 #' @param prob Probabilities of gang demand corresponding to "gangs."
12 #'

```



```

13 #' @return Number representing the total gang demand
14 demand_gangs <- function(n_ships, gangs = gang, prob = p) {
15   if (length(gangs) == 1) {
16     gangs <- c(gangs, gangs)
17     prob <- c(1, 1)
18   }
19   requests <- sample(
20     gangs,
21     size = n_ships,
22     replace = TRUE,
23     prob = prob
24   )
25   sum(requests)
26 }

```

# Using the simulation function

- As an example, we can simulate the **total gang request** for  $s = 10$  ships:

R Code ↺ Start Over ▶ Run Code

```
1 set.seed(551) # Reproducibility
2 demand_gangs(10)
```

## If the number of ships arriving is random

- Suppose  $S$ , the number of ships arriving on a given day, follows a **Poisson distribution with a mean of  $\lambda = 5$  ships**.
- What is the distribution of total gang requests on a given day?

## Two-step simulation

1. Generate the number of ships  $S \sim \text{Poisson}(\lambda = 5 \text{ ships})$  for  $n$  days .
2. For each day, simulate the total gang request  $T = \sum_{i=1}^S D_i$  for the simulated number of ships.
3. We now have a random sample of size  $n$ .

# The Code

- Let us try this, obtaining a sample of  $n = 10000$  days.

R Code [↺ Start Over](#) [▶ Run Code](#)

```
1 n_days <- 10000
2
3 # Setting global seed!
4 set.seed(551)
5
6 ## Step 1: generate a bunch of ships arriving each day.
7 arrivals <- rpois(n_days, lambda = 5)
8 head(arrivals)
```

R Code [↺ Start Over](#) [▶ Run Code](#)

```
1 ## Step 2: Simulate the grand total gang request on each day.
2 library(purrr)
3 total_requests <- map_int(arrivals, demand_gangs)
4 head(total_requests)
```

# Then...

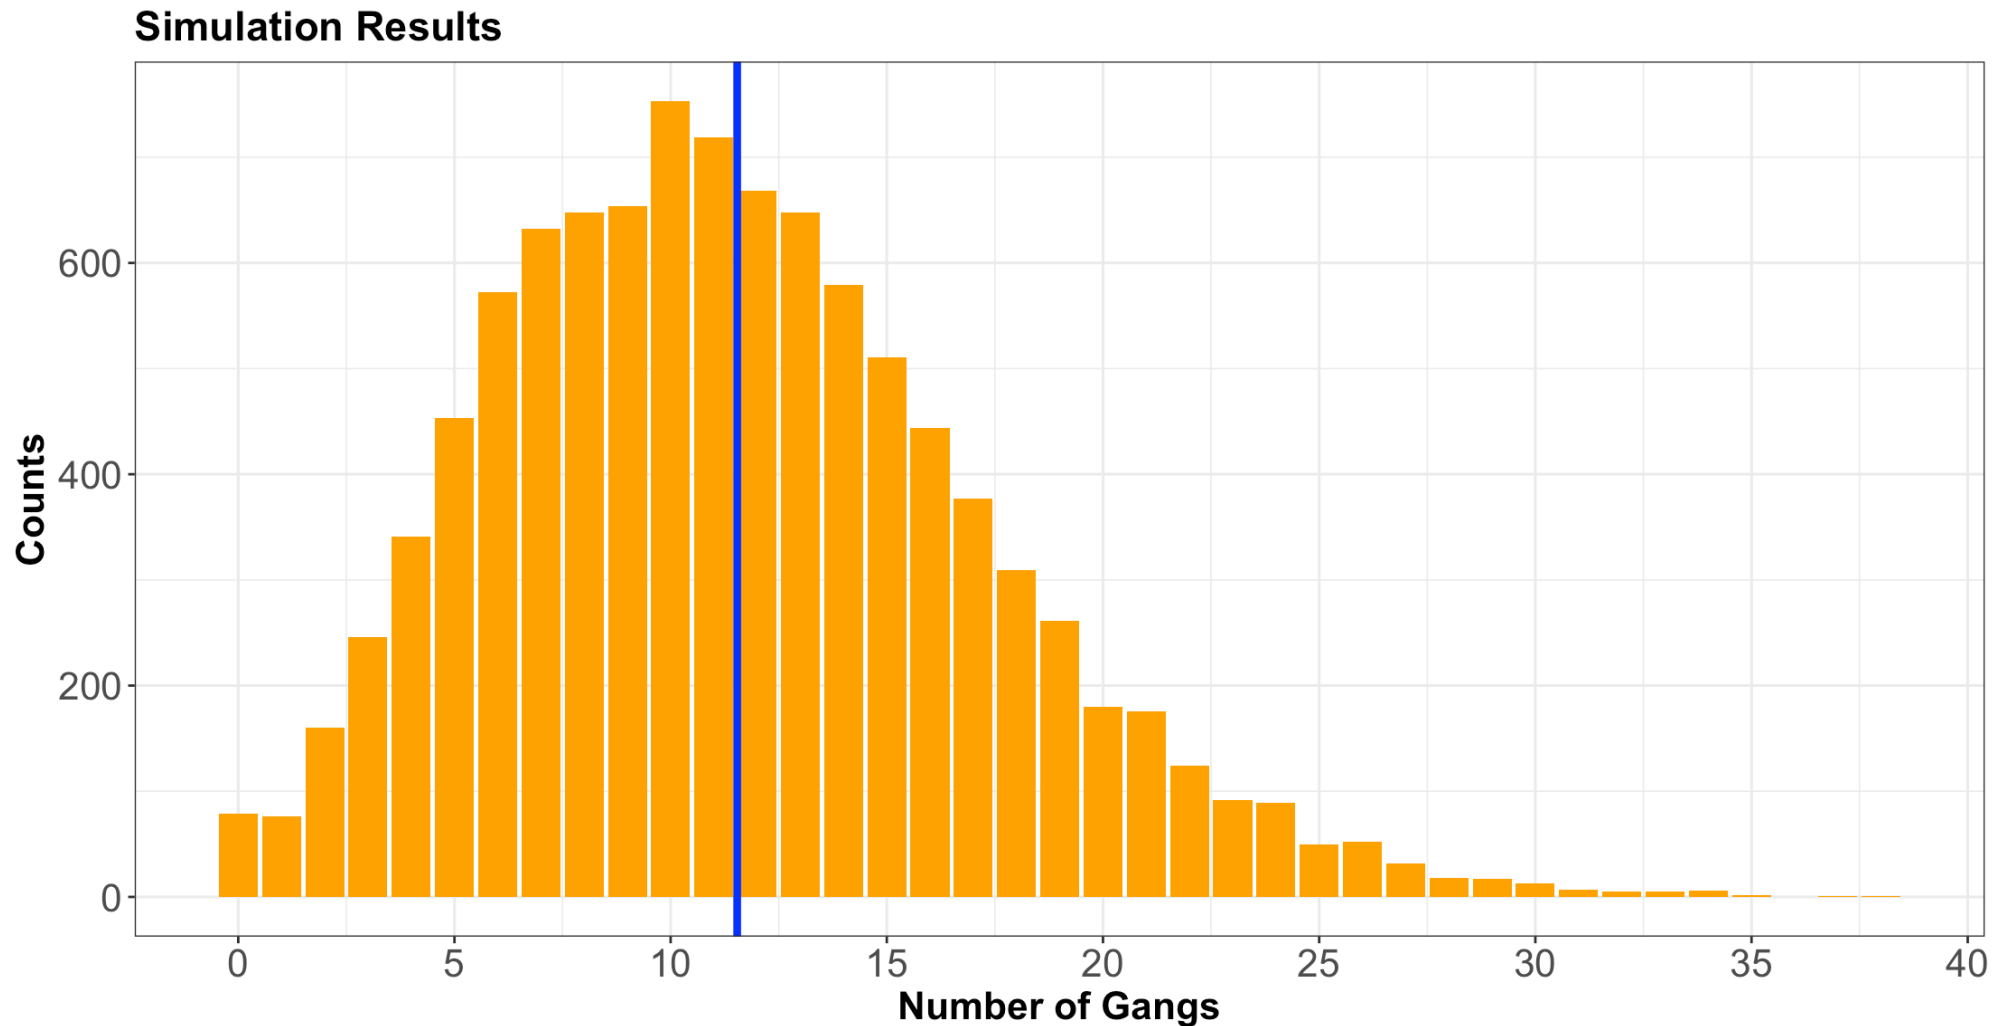
R Code

[↺ Start Over](#)[▶ Run Code](#)

```
1 ## Step 3: Compute mean and variance.
2 simulation_outputs <- tibble(
3   mean = mean(total_requests),
4   variance = var(total_requests)
5 )
6 simulation_outputs
```

- For  $n = 10000$ , the summary statistics indicate that **we expect a demand of 11.5 gangs on a given day in the whole port.**

# Checking the empirical distribution of total gang requests!



# Today's Learning Goals

- Generate a random sample from a distribution in **R**.
- Reproduce the same random sample each time you re-run your code in **R** by setting the seed or random state.
- Use simulation to approximate distribution properties (e.g., mean and variance), especially for random variables involving multiple other random variables.
- Understand why simulations can approximate true properties of a distribution.



# Final Wrap Up

- Probability is everywhere: finance, computer science, physics, engineering, and data science!
- You will learn more advanced topics that build on what we've covered, such as regression, Bayesian inference, machine learning models, and LLMs.